



Cassandra

Cassandra


By: Sheefali Tewari

What is Cassandra?

- Scaled NoSQL database
- Good for managing large amounts of structured, semi-structured, and unstructured data
- No single point of failure

Features

- Elastic scalability
- Always on architecture
- Fast linear-scale performance
 - Example:



More Features

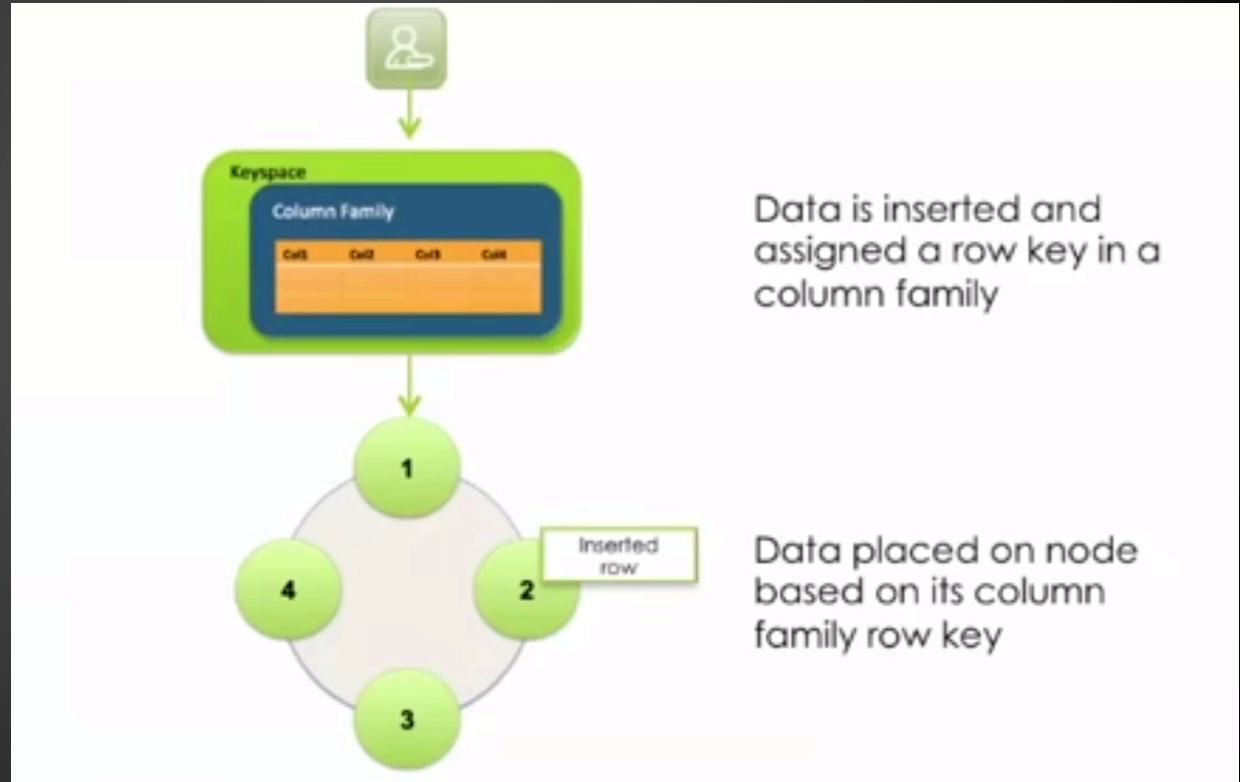
- Flexible data storage
- Easy data distribution
- Operational simplicity

Even More Features

- Transaction support
- Availability

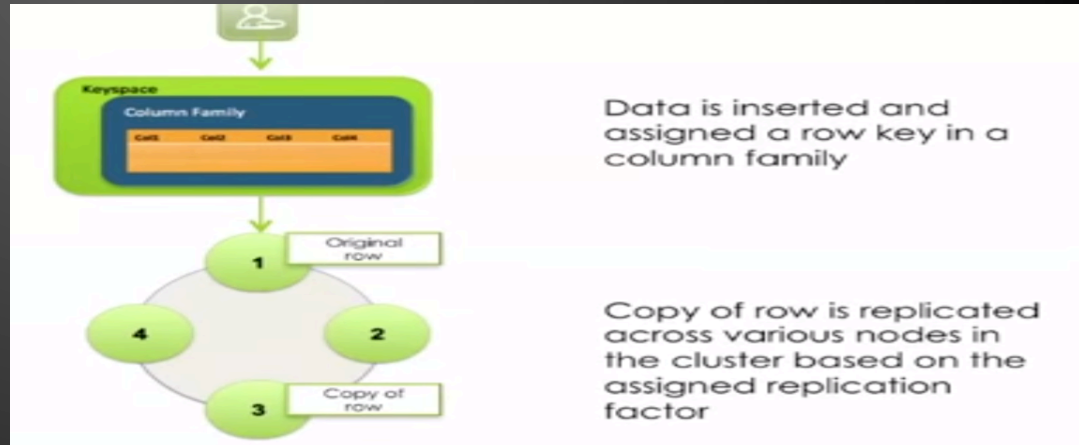
Data Partitioning in Cassandra

- Random
- Ordered

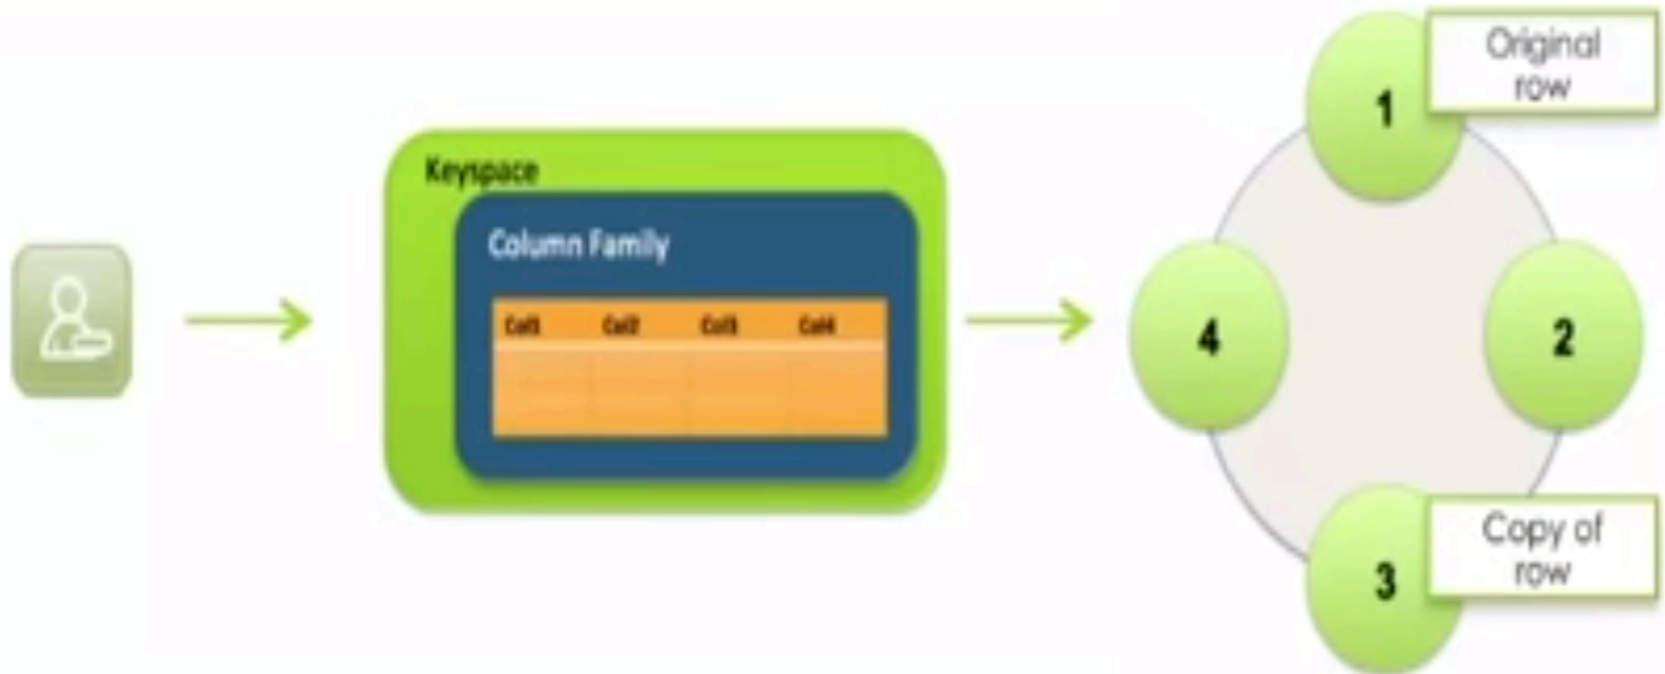


Replication in Cassandra

- To ensure fault tolerance
- No single point of failure

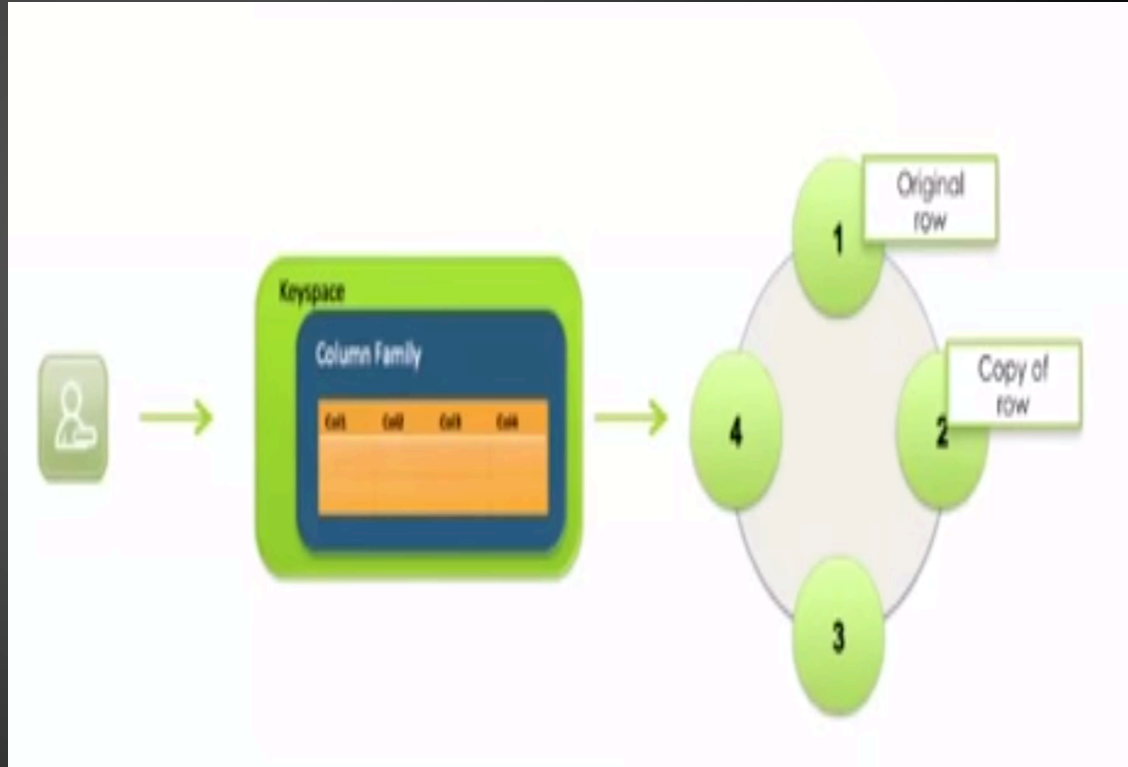


Replication in Cassandra

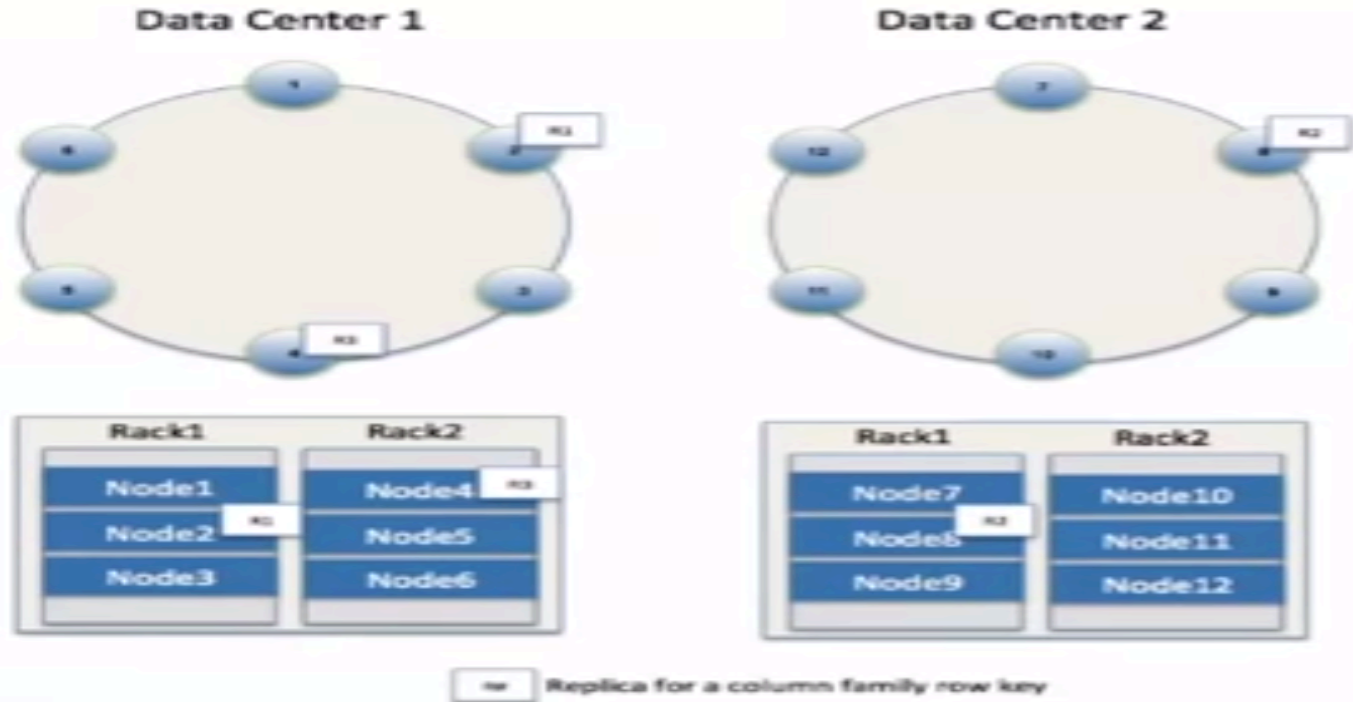


Replication Strategies: Simple

- Place original row on a node determined by the partitioner



Replication Strategies: Network Topology Strategy



Comparisons

- Apache Cassandra: scalable, high performance distributed database
- Apache HBase: open source, non-relational, distributed database
- MongoDB: cross-platform document-oriented database system

Who Uses Cassandra?

- Netflix
- Ebay
- Twitter
- Reddit
- Etc

Source

<http://planetcassandra.org/what-is-apache-cassandra/>