

Riak

-Priyanka Goyal



What is Riak?

- Distributed Key-value database
 - ◆ bucket
 - key-value
 - key
 - value
- Can store plain text, JSON, XML, images, videos
- Fault-tolerant
- Simple REST API
- Best suitable for Amazon

Features

- **SCALING:** Value of key defines on which node key is stored
- **CONSISTENCY:** Eventual but how to resolve update conflicts?
 - ◆ 2 ways
 - newest write wins, older write loose
 - return both values to client and get it resolved
- **TRANSACTIONS:** Uses Quorum concept

Quorums - The Riak Way!

- n = The number of Replicas (Replication Factor)
 - varies per bucket
- r = The number of Replicas needed for successful read
 - varies per operation
- w = The number of Replicas needed for successful write
 - varies per operation

Quorums - The Riak Way! (Contd..)



$n-r$ = read fault tolerance

$n-w$ = write fault tolerance

Example: 6-node Riak cluster, $n=4, r=1, w=2$

- $4-1 = 3$ hosts can be down but Riak will still perform read operation
- $4-2 = 2$ hosts can be down but Riak will still perform write operation

Map/Reduce

Map() : Convert list of data into another type of list

Reduce() : convert the new list to one/more scalar values

The Riak Way!

Each object on the server is “mapped” to a common key that groups data together and then all matching keys are “reduced” into a single value.

Erlang Map Phase

```
map_object_value(obj, _keyData, _Arg)->[riak_object:  
    get_value(obj)]
```

- obj: Riak object retrieved from bucket/key
- keyData: Static argument specified with the bucket/key
- Arg: Static argument specified with the job

Built-in Functions:

- riak_mapreduce:map_object_value/3
- riak_mapreduce:map_object_value_list/3

Erlang Reduce Phase

Built-in Functions:

- `riak_mapreduce:reduce_set_union/2`
 - `riak_mapreduce:reduce_sum/2`
 - `riak_mapreduce:reduce_sort/2`
-
- Riak explores MapReduce using REST API
 - Riak submits job via POST
 - Default URL is `/mapred`



- ```
</riak/demo/test1>;riaktag="userinfo"
```

- Examples:** /riak/demo/test1/\_,\_,1

```
/riak/demo/test1/demo,_,1
```

```
/riak/demo/test1/_,_,0/_,_,1
```

```
/riak/demo/test1/_,child,0/_,_,1
```

## When to use

- Storing Session Information
  - ◆ web-sessions with their unique sessionId. Use PUT/GET
- User Profiles, Preferences
- Shopping Cart Data
  - ◆ e-commerce websites

## When not to use

- Relationship among Data
  - ◆ correlate data b/w different sets of keys
- Multiops Transactions
  - ◆ multiple keys with failure to save even one key
- Query by Data
  - ◆ search keys based on values
- Operations on multiple keys

# Wrap up!

...flexible storage engine...  
...with REST API...  
...and Map/Reduce capability...  
...designed to be fault-tolerant...  
...distributed...  
...operations friendly...