



Your very own search engine!

# What's this Lucene thing?

- Apache Lucene is a high-performance, full-featured text search engine library written entirely in Java. It is a technology suitable for nearly any application that requires full-text search, especially cross-platform ones.
- This type of index is called an inverted index, because it inverts a page-centric data structure (page->words) to a keyword-centric data structure (word->pages).
- **Important:** instead of searching the text directly, Lucene searches an index instead.




# What is Solr?



- At it's most basic it's a search server built on top off Apache Lucene that provides an array of features.
- On the website: “highly reliable, scalable and fault tolerant, providing distributed indexing, replication and load-balanced querying, automated failover and recovery, centralized configuration and more”
- How I've used it: a simple and high performance way to do database searches
- Reference manual is over 500 pages.

# How does Solr work?

- A. Solr Searches the Indexed Documents.
- B. Each Document has an ID and list of terms
- C. For Each term, solr has lists of all document that contains this specific term.
- D. Identify each document by it's ID and return it.



TL;DR - Solr is a Lucene  
Frontend that provides  
fast text search on a DB.

Now onto the  
fun stuff!



# Lets see it in action...

(A sacrifice to the live  
action Gods)



## A real use case:

CU FCQ - A Ruby on Rails project to visualize CU's FCQs.

- Dirty scraped data.
- Maintain a DB with over 100,000 rows.
- Needs to be searchable.
- Needs to be fast.



# What I was doing before:

```
def self.search(query)
  #If there wasn't any input return everything.
  if query.nil?
    return Instructors.all
  end
  #this makes capital and strips whitespace
  q = query.upcase.strip
  #actually make an sql query string using the input
  if q.include?(" ")
    where("instructor_last like ? AND instructor_first like ? OR instructor_first
like ? AND instructor_last like ?",q.split[1],q.split[0],q.split[1],q.split[0])
  else
    where("instructor_last = ? OR instructor_first = ? ",q,q)
  end
end
```

## Notes:


1. Don't laugh.
2. This was in my controller.
3. I was taking input directly from an HTML input box.
4. Use the input to create a SQL string and search.

This is bad on a number of levels.



# Enter Sunspot

- Where to begin: I Googled “Search Engine in Rails” and came up with this [gem](#).
  - A gem is a little library that you can download to plugin into your app.
- Sunspot is built on top of the RSolr library, which provides a low-level interface for Solr interaction; Sunspot provides a simple, intuitive, expressive DSL backed by powerful features for indexing objects and searching for them.
- **TL;DR Comes with a bundled version of Solr, so boilerplate is good to go out of the box.**



# Integrating it into my project (pt. 1)

## Models (Our Logic):

- Add the searchable fields for the classes.

Instructor.rb-

```
searchable do
  text :instructor_first
  text :instructor_last, :default_boost => 2
end
```

Course.rb-

```
searchable do
  text :crse
  text :subject
  text :course_title, :default_boost => 5
end
```

## Controller (DB <-> View):

- When we are calling an index for instructors we can perform our search.
- We also get pagination (30 by default)

def index

```
@search = Instructor.search do
```

```
  fulltext params[:search]
```

```
  paginate :page => params[:page] || 1
```

```
end
```

```
@instructors = @search.results
```

```
end
```

# Displaying it to user (pt. 2)

## Our Instructors Index view:

- This is where we list the return values of our query

```
<th>Instructor</th>
<th>Instructor Type</th>
...
<th>Last FCQ</th>

<tbody>
  <% @instructors.each do |instructor| %>
    <td><%= link_to instructor.name, instructor %>
  ></td>
    <td><%= instructor.instr_group %></td>

    ....
    <td><%= instructor.latest_teaching %></td>
  </tbody>
  <% end %>
<%= will_paginate @instructors %>
```

## The Search Box :

- We add the search box so the user can interact with Solr.
- Anything with `<%= ... %>` is Ruby code.
- Note: I removed indentation so it's more readable.

```
<form class="navbar-form navbar-right">
  <%= form_tag(instructors_path, :method => "get", id: "form-inline")
do %>
  <%= text_field_tag :search, params[:search], placeholder: "Search"
%>

  <%= submit_tag "Search", name: nil, :class=>'button' %>
<% end %>
```

# Putting it all together

- At this point, simply populate your DB again. (Pitfall 0)
- Start Solr and reindex everything
  - Note: Every time you need to change what Solr is indexing (new terms, new weights, etc.), you need to reindex.
- Start your web server and search!



# Pitfalls

# Pitfall 0: Solr and Rake

- To initialize our DB, we need to make sure Solr is killed.
- We want to remove all of our old temporary information and indices.
- Before we start actually putting Data into the DB, we restart our Solr server.
- Run our rake tasks.
- Reindex Solr once our DB is full.

```
#!/bin/bash

#kill solr process
pkill -f solr

#remove old data
rm -rf solr/data
rm -rf solr/default
rm -rf solr/development
rm -rf solr/pids
rm -rf solr/test

#startup solr in development environment
rake sunspot:solr:start RAILS_ENV=development
#all of our rake tasks (yours will be different)
bundle exec rake db:reset

...

bundle exec rake grades

#reindex solr
RAILS_ENV=development bundle exec rake sunspot:solr:
reindex
```



# Pitfall 1: Starting your Server w/ Solr

- Just starting Sunspot can be a challenge.
- You need to specify what type of server you are going to be using with Sunspot (development)
- Each of these listen on different ports, so you might get “Connection Refused” or “404” errors if Sunspot can’t connect to Solr.
- This script starts and reindexes Solr along with our web server.

```
#!/bin/bash
```

```
echo "Starting/Reindexing Solr"
```

```
rake sunspot:solr:start RAILS_ENV=development
```

```
rake sunspot:solr:reindex RAILS_ENV=development
```

```
echo "Solr successful! Starting Rails"
```

```
rails server -b 0.0.0.0 -p 80
```

```
echo "rails server closed!"
```



# Pitfall 2: Solr configs and Git


- You want to be able share your project on Git.
- The solr/ directory contains import config information and temporary data (pid,index,etc.)
- Add the following code to your .gitignore so you don't start tracking any temp data that could wreck clones.
- You are still tracking import config [information](#).
- Only keep conf/ and solr.xml

```
#Ignore all of the solr stuff except the xml and the  
conf  
  
solr/data  
solr/default  
solr/development  
solr/pids  
solr/test
```

# Conclusion

- Solr is a super cool and (relatively) simple way to get search functionality for your DB.
- It's fast, reliable, and has fantastic features like pagination and indexing.
- Interested in CUFCQ?
  - This is our Alpha launch and we're looking for help.

You can access us @ [cufcq.com](http://cufcq.com)



# Important/Interesting Links:

- Integrating Rails with Solr [tutorial](#).
- Basic [explanation](#) of Solr.
- Solr [Sunspot](#).
- Our [Project](#).

Need help? Email me (Alex Tsankov) at:  
[antsankov@gmail.com](mailto:antsankov@gmail.com)